

Tuning a PID Controller

These steps apply to position PID controllers. Velocity PID controllers typically don't need .

1. Set K_p , K_i , and K_d to zero.
2. Increase K_p until the [output](#) starts to oscillate around the [setpoint](#).
3. Increase K_p as much as possible without introducing jittering in the [system response](#).

Plot the position [setpoint](#), velocity [setpoint](#), measured position, and measured velocity. The velocity [setpoint](#) can be obtained via numerical differentiation of the position [setpoint](#) (i.e., \dot{x}). Increase K_p until the position tracks well, then increase until the velocity tracks well.

If the [controller](#) settles at an [output](#) above or below the [setpoint](#), one can increase K_i such that the [controller](#) reaches the [setpoint](#) in a reasonable amount of time. However, a steady-state feedforward is strongly preferred over integral control (especially for PID control).

Important

Adding an integral gain to the [controller](#) is an incorrect way to eliminate [steady-state error](#). A better approach would be to tune it with an integrator added to the [plant](#), but this requires a [model](#). Since we are doing output-based rather than model-based control, our only option is to add an integrator to the [controller](#).

Beware that if K_i is too large, integral windup can occur. Following a large change in [setpoint](#), the integral term can accumulate an error larger than the maximal [control input](#). As a result, the system overshoots and continues to increase until this accumulated error is unwound.

Note

The [frc-characterization toolsuite](#) can be used to model your system and give accurate Proportional and Derivative values. This is preferred over tuning the controller yourself.

Actuator Saturation¶

A controller calculates its output based on the error between the [reference](#) and the current [state](#). [Plant](#) in the real world don't have unlimited control authority available for the controller to apply. When the actuator limits are reached, the controller acts as if the gain has been temporarily reduced.

We'll try to explain this through a bit of math. Let's say we have a controller where u is the [control effort](#), K_p is the gain, r is the [reference](#), and x is the current state. Let u_{max} be the limit of the actuator's output which is less than the uncapped value of u and be the

associated maximum gain. We will now compare the capped and uncapped controllers for the same [reference](#) and current [state](#).

For the inequality to hold, must be less than the original value for . This reduced gain is evident in a [system response](#) when there is a linear change in state instead of an exponential one as it approaches the [reference](#). This is due to the [control effort](#) no longer following a decaying exponential plot. Once the [system](#) is closer to the [reference](#), the controller will stop saturating and produce realistic controller values again.